

## CSE278: Introduction to Systems Programming (Systems I)

### Lab #9

**Due: Mon/Tue May 4/5 during Lab time**

**Maximum Points: 50**

#### Submission Instructions

This part of the homework assignment must be turned-in electronically via Canvas. Ensure you name this document `Lab9_MUID.docx`, where `MUID` is your Miami University unique ID. (Example: `Lab9_ahmede.docx`)

Copy pasting from online resources is **Plagiarism**. Instead you should read, understand, and use your own words to respond to questions.

#### Submission Instructions:

Once you have completed answering the questions save this document as a PDF file (**don't just rename the document; that is not the correct way to save as PDF**) and upload it to Canvas.

**General Note:** Upload each file associated with homework (or lab exercises) individually to Canvas. Do not upload archive file formats such as zip/tar/gz/7zip/rar etc.

#### Objective

The objective of this Lab is to review basic concepts of:

- Linux ABI
- Debugging with gdb

Name: **Ben Hilger**

### Required reading

- Lecture Slides & ClassNotes: Security
- Lecture Slides & ClassNotes: ComputerArchitecture

### PART A: Linux Application Binary Interface (ABI)

#### Goals

- Details of how a program executes on a Unix/Linux system

**Linux ABI defines most of the low-level details of a program including:**

- The register layout (rip, rsp, rbp, rax,, rbx, rcx, rdx, rdi, rsi, r8, r9, r10, r11, r12, r13, r14, and r15)
- The stack frame
  - Pushing to the stack *subtracts* from the stack pointer
  - Popping from the stack *adds* to the stack pointer
  - call
  - ret
- Function prologs and epilogs
- The calling convention (that is, parameter passing)
- Exception handling
- Virtual memory layout
- The binary object format (ELF) (begins with 0x7F)

1. Consider the following code fragment:

```
int test()
{
    int i = 1;
    int j = 2;
    return i + j;
}
int main (void)
{
    test();
}
```

A. Compile the Lab9.cpp as follows:

**g++ Lab9.cpp**

B. Disassemble the resulting binary using the command

**objdump -S a.out**

C. **Attach a screen shot**

```
hilgerbj@os1:~/cse278/Lab9$ touch Lab9.cpp
hilgerbj@os1:~/cse278/Lab9$ nano Lab9.cpp
hilgerbj@os1:~/cse278/Lab9$ g++ Lab9.cpp
hilgerbj@os1:~/cse278/Lab9$ objdump -S a.out

a.out:      file format elf64-x86-64

Disassembly of section .init:

0000000000004b0 <.init>:
4b8:  48 83 ec 08          sub    $0x8,%rsp
4bc:  48 8b 05 25 0b 20 00 mov    0x200b25(%rip),%rax      # 200fe8 <__gmon_start__>
4c3:  48 85 c0             test   %rax,%rax
4c6:  74 02              je     4ca <.init+0x12>
4c8:  ff d0             callq  *%rax
4ca:  48 83 c4 08          add    $0x8,%rsp
4ce:  c3                retq

Disassembly of section .plt:

0000000000004d0 <.plt>:
4d0:  ff 35 f2 0a 20 00    pushq 0x200af2(%rip)          # 200fc8 <_GLOBAL_OFFSET_TABLE_+0x8>
4d6:  ff 25 f4 0a 20 00    jmpq   *0x200af4(%rip)        # 200fd0 <_GLOBAL_OFFSET_TABLE_+0x10>
4dc:  0f 1f 40 00          nopl   0x0(%rax)

Disassembly of section .plt.got:

0000000000004e0 <__cxa_finalize@plt>:
4e0:  ff 25 12 0b 20 00    jmpq   *0x200b12(%rip)        # 200ff8 <__cxa_finalize@GLIBC_2.2.5>
4e6:  66 90              xchgb  %ax,%ax

Disassembly of section .text:

0000000000004f0 <_start>:
4f0:  31 ed              xor     %ebp,%ebp
4f2:  49 89 d1            mov     %rdx,%r9
4f5:  5e                pop     %rsi
4f6:  48 89 e2            mov     %rsp,%rdx
4f9:  48 83 e4 f0          and     $0xfffffffffffffff0,%rsp
4fd:  50                push    %rax
4fe:  54                push    %rsp
4ff:  4c 8d 05 9a 01 00 00 lea     0x19a(%rip),%r8        # 6a0 <__libc_csu_fini>
506:  48 8d 0d 23 01 00 00 lea     0x123(%rip),%rcx       # 630 <__libc_csu_init>
50d:  48 8d 3d 02 01 00 00 lea     0x102(%rip),%rdi       # 616 <main>
514:  ff 15 c6 0a 20 00    callq  *0x200ac6(%rip)        # 200fe0 <__libc_start_main@GLIBC_2.2.5>
51a:  f4                hlt
51b:  0f 1f 44 00 00       nopl   0x0(%rax,%rax,1)

000000000000520 <deregister_tm_clones>:
520:  48 8d 3d e9 0a 20 00 lea     0x200ae9(%rip),%rdi     # 201010 <__TMC_END__>
527:  55                push    %rbp
528:  48 8d 05 e1 0a 20 00 lea     0x200ae1(%rip),%rax     # 201010 <__TMC_END__>
52f:  48 39 f8            cmp     %rdi,%rax
532:  48 89 e5            mov     %rsp,%rbp
535:  74 19              je      550 <deregister_tm_clones+0x30>
537:  48 8b 05 9a 0a 20 00 mov     0x200a9a(%rip),%rax     # 200fd8 <_ITM_deregisterTMCloneTable>
53e:  48 85 c0            test    %rax,%rax
541:  74 0d              je      550 <deregister_tm_clones+0x30>
543:  5d                pop     %rbp
544:  ff e0             jmpq    *%rax
546:  66 2e 0f 1f 84 00 00 nopw    %cs:0x0(%rax,%rax,1)
54d:  00 00 00
550:  5d                pop     %rbp
551:  c3                retq
552:  0f 1f 40 00       nopl   0x0(%rax)
556:  66 2e 0f 1f 84 00 00 nopw    %cs:0x0(%rax,%rax,1)
55d:  00 00 00
```

```

000000000000560 <register_tm_clones>:
560: 48 8d 3d a9 0a 20 00 lea 0x200aa9(Nrip),Nrdi # 201010 <__TMC_END__>
567: 48 8d 35 a2 0a 20 00 lea 0x200aa2(Nrip),Nrsl # 201010 <__TMC_END__>
56e: 55 push Nrbp
56f: 48 29 fe sub Nrdi,Nrsl
572: 48 89 e5 mov Nrsp,Nrbp
575: 48 c1 fe 03 sar $0x3,Nrsl
579: 48 89 f0 mov Nrsl,Nrax
57c: 48 c1 e8 3f shr $0x3f,Nrax
580: 48 01 c6 add Nrax,Nrsl
583: 48 d1 fe sar Nrsl
586: 74 18 je 5a0 <register_tm_clones+0x40>
588: 48 0b 05 61 0a 20 00 mov 0x200aa1(Nrip),Nrax # 200ff0 <__ITM_registerTMCloneTable>
58f: 48 85 c0 test Nrax,Nrax
592: 74 0c je 5a0 <register_tm_clones+0x40>
594: 5d pop Nrbp
595: ff e0 jmpq *Nrax
597: 66 0f 1f 84 00 00 00 nopw 0x0(Nrax,Nrax,1)
59e: 00 00
5a0: 5d pop Nrbp
5a1: c3 retq
5a2: 0f 1f 40 00 nopl 0x0(Nrax)
5a5: 66 2e 0f 1f 84 00 00 nopw Ncs:0x0(Nrax,Nrax,1)
5ad: 00 00 00

0000000000005b0 <__do_global_ctors_aux>:
5b0: 80 3d 59 0a 20 00 00 cmpb $0x0,0x200aa9(Nrip) # 201010 <__TMC_END__>
5b7: 75 2f jne 5e8 <__do_global_ctors_aux+0x38>
5b9: 48 83 3d 37 0a 20 00 cmpq $0x0,0x200aa37(Nrip) # 200ff8 <__cxa_finalize@GLIBC_2.2.5>
5c0: 00
5c1: 55 push Nrbp
5c2: 48 89 e5 mov Nrsp,Nrbp
5c5: 74 0c je 5d3 <__do_global_ctors_aux+0x23>
5c7: 48 0b 3d 3a 0a 20 00 mov 0x200aa3a(Nrip),Nrdi # 201008 <__dso_handle>
5ce: e8 0d ff ff ff callq 4e0 <__cxa_finalize@plt>
5d3: e8 48 ff ff ff callq 520 <deregister_tm_clones>
5d8: c6 05 31 0a 20 00 01 movb $0x1,0x200aa31(Nrip) # 201010 <__TMC_END__>
5df: 5d pop Nrbp
5e0: c3 retq
5e1: 0f 1f 80 00 00 00 00 nopl 0x0(Nrax)
5e8: f3 c3 repz retq
5ea: 66 0f 1f 44 00 00 00 nopw 0x0(Nrax,Nrax,1)

0000000000005f0 <frame_dummy>:
5f0: 55 push Nrbp
5f1: 48 89 e5 mov Nrsp,Nrbp
5f4: 5d pop Nrbp
5f5: e9 66 ff ff ff jmpq 560 <register_tm_clones>

0000000000005fa <_Z4testv>:
5fa: 55 push Nrbp
5fb: 48 89 e5 mov Nrsp,Nrbp
5fe: c7 45 f8 01 00 00 00 movl $0x1,-0x0(Nrbp)
605: c7 45 fc 02 00 00 00 movl $0x2,-0x4(Nrbp)
60c: 8b 55 f8 mov -0x0(Nrbp),Nedx
60f: 8b 45 fc mov -0x4(Nrbp),Neax
612: 01 d0 add Nedx,Neax
614: 5d pop Nrbp
615: c3 retq

000000000000616 <main>:
616: 55 push Nrbp
617: 48 89 e5 mov Nrsp,Nrbp
61a: e8 0b ff ff ff callq 5fa <_Z4testv>
61f: b8 00 00 00 00 mov $0x0,Neax
624: 5d pop Nrbp
625: c3 retq
626: 66 2e 0f 1f 84 00 00 00 nopw Ncs:0x0(Nrax,Nrax,1)
62d: 00 00 00

1000000000000630 <__libc_csu_init>:
630: 41 57 push Nrl5
632: 41 56 push Nrl4
634: 49 89 d7 mov Nrdx,Nrl5
637: 41 55 push Nrl3
639: 41 54 push Nrl2
63b: 4c 8d 25 ae 07 20 00 lea 0x2007ae(Nrip),Nrl2 # 200df0 <__frame_dummy_init_array_entry>
642: 55 push Nrbp
643: 48 8d 2d ae 07 20 00 lea 0x2007ae(Nrip),Nrbp # 200df8 <__init_array_end>
64a: 53 push Nrbx
64b: 41 89 fd mov Nedi,Nrl3d
64d: 49 89 f6 mov Nrsl,Nrl4
651: 4c 29 e5 sub Nrl2,Nrbp
654: 48 83 ec 08 sub $0x8,Nrsp
658: 48 c1 fd 03 sar $0x3,Nrbp
65c: e8 07 fe ff ff callq 480 <_init>
661: 48 85 ed test Nrbp,Nrbp
664: 74 20 je 686 <__libc_csu_init+0x56>
666: 31 db xor Nedx,Nedx
668: 0f 1f 84 00 00 00 00 nopl 0x0(Nrax,Nrax,1)
66f: 00
670: 4c 89 fa mov Nrl5,Nrdx
673: 4c 89 f6 mov Nrl4,Nrsl
676: 44 89 ef mov Nrl3d,Nedi
679: 42 ff 14 dc callq *(Nrl2,Nrbx,8)
67d: 48 83 c3 01 add $0x1,Nrbx
681: 48 39 dd cmp Nrbx,Nrbp
684: 75 ea jne 670 <__libc_csu_init+0x48>
686: 48 83 c4 08 add $0x8,Nrsp
68a: 5b pop Nrbx
68d: 5d pop Nrbp
68e: 41 5c pop Nrl2
68e: 41 5d pop Nrl3
690: 41 5e pop Nrl4
692: 41 5f pop Nrl5
694: c3 retq
695: 90 nop
696: 66 2e 0f 1f 84 00 00 00 nopw Ncs:0x0(Nrax,Nrax,1)
69d: 00 00 00

10000000000006a0 <__libc_csu_fini>:
6a0: f3 c3 repz retq

Disassembly of section .fini:
10000000000006a4 <_fini>:
6a4: 48 83 ec 08 sub $0x8,Nrsp
6a8: 48 83 c4 08 add $0x8,Nrsp
6ac: c3 retq

```

2. Now modify the code as follows:

```

int test( int val1, int val2)
{
    return val1 + val2;
}

int main (void)
{
    auto ret = test (42, 42);
}

```

```
}
```

- a. Again from the resulting binary, report the generated assembly language code for relevant for the **main()** function and **test()** function
- b. Issue the following command:

**readelf -SW a.out**

How many sections reported there?

There are 28 sections reported

- c. To know more details of readelf, issue the command  
**man readelf**
- d. Issue the following command:  
**readelf --debug-dump=frames a.out**

**Report the contents of the .eh\_frame table**

```
hligert@os1:~/cse278/Lab9$ readelf --debug-dump=frames a.out
Contents of the .eh_frame section:

00000000 0000000000000014 00000000 CIE
Version: 1
Augmentation: "zR"
Code alignment factor: 1
Data alignment factor: -8
Return address column: 16
Augmentation data: 1b
DW_CFA_def_cfa: r7 (rsp) ofs 8
DW_CFA_offset: r16 (rip) at cfa-8
DW_CFA_undefined: r16 (rip)

00000018 0000000000000014 0000001c FDE cie=00000000 pc=00000000000006f0..000000000000071b
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop

00000030 0000000000000014 00000000 CIE
Version: 1
Augmentation: "zR"
Code alignment factor: 1
Data alignment factor: -8
Return address column: 16
Augmentation data: 1b
DW_CFA_def_cfa: r7 (rsp) ofs 8
DW_CFA_offset: r16 (rip) at cfa-8
DW_CFA_nop
DW_CFA_nop

00000048 0000000000000024 0000001c FDE cie=00000030 pc=00000000000006a0..00000000000006e0
DW_CFA_def_cfa_offset: 16
DW_CFA_advance_loc: 6 to 00000000000006a6
DW_CFA_def_cfa_offset: 24
DW_CFA_advance_loc: 10 to 00000000000006b0
DW_CFA_def_cfa_expression (DW_OP_breg7 (rsp): 8; DW_OP_breg16 (rip): 0; DW_OP_lit15; DW_OP_and; DW_OP_lit11; DW_OP_ge; DW_OP_lit3; DW_OP_shl; DW_OP_plus)
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop

00000070 0000000000000014 00000044 FDE cie=00000030 pc=00000000000006e0..00000000000006e8
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
DW_CFA_nop
```

The .eh\_frame contains sections of CIE (Common Information Entry) and FDE (Frame Description Entry) that handle certain parts of exceptions. Every CIE is followed by one or more FDE in the output and each memory address comes one after the other.

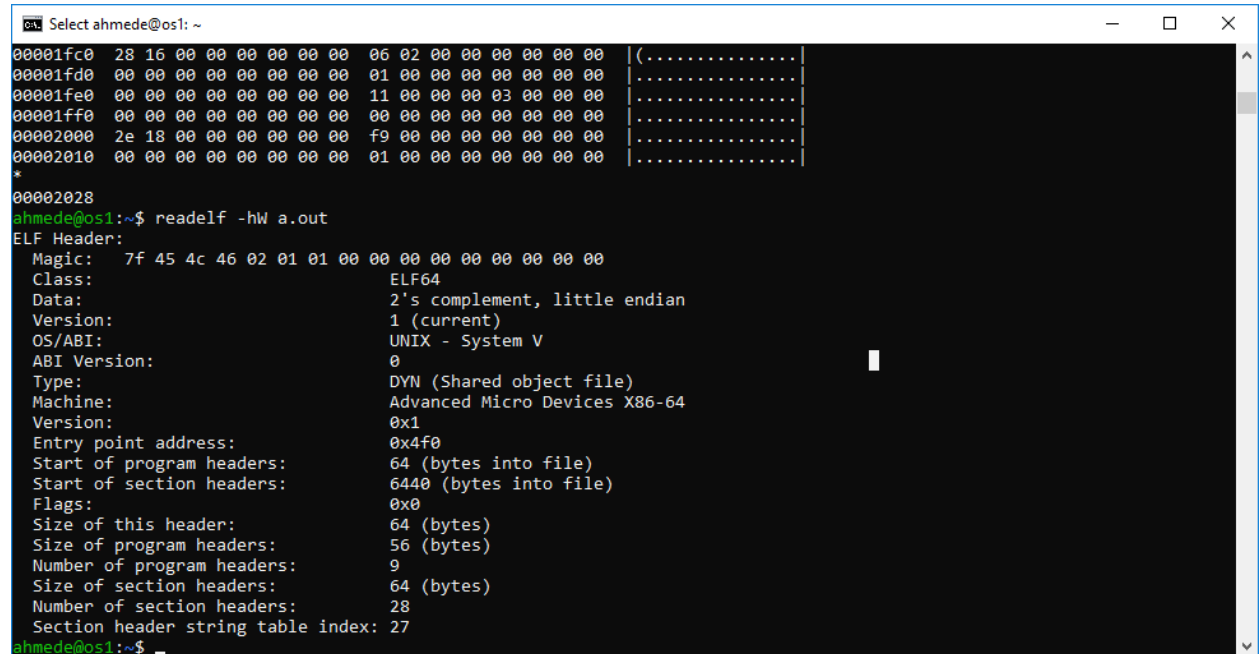
- e. Look at the *hexdump* of the resulting a.out ELF by issuing the following command:  
**hexdump -C a.out**

Every ELF file begins with the hex number 0x7F, and continues with the ELF string.

Issue the following command to view the ELF file's header:

**readelf -hW a.out**

You should see a result like this:



```
Select ahmede@os1: ~
00001fc0 28 16 00 00 00 00 00 00 06 02 00 00 00 00 00 00 |(.|.....|
00001fd0 00 00 00 00 00 00 00 00 01 00 00 00 00 00 00 00 |.....|
00001fe0 00 00 00 00 00 00 00 00 11 00 00 00 03 00 00 00 |.....|
00001ff0 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
00002000 2e 18 00 00 00 00 00 00 f9 00 00 00 00 00 00 00 |.....|
00002010 00 00 00 00 00 00 00 00 01 00 00 00 00 00 00 00 |.....|
*
00002028
ahmede@os1:~$ readelf -hW a.out
ELF Header:
  Magic:      7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:      ELF64
  Data:       2's complement, little endian
  Version:    1 (current)
  OS/ABI:     UNIX - System V
  ABI Version: 0
  Type:       DYN (Shared object file)
  Machine:    Advanced Micro Devices X86-64
  Version:    0x1
  Entry point address: 0x4f0
  Start of program headers: 64 (bytes into file)
  Start of section headers: 6440 (bytes into file)
  Flags:      0x0
  Size of this header: 64 (bytes)
  Size of program headers: 56 (bytes)
  Number of program headers: 9
  Size of section headers: 64 (bytes)
  Number of section headers: 28
  Section header string table index: 27
ahmede@os1:~$
```

Attach the screen shot what you get from previous command.

```
hilgerbj@os1:~/cse278/Lab9$ readelf -hW a.out
ELF Header:
  Magic:      7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:      ELF64
  Data:       2's complement, little endian
  Version:    1 (current)
  OS/ABI:     UNIX - System V
  ABI Version: 0
  Type:       DYN (Shared object file)
  Machine:    Advanced Micro Devices X86-64
  Version:    0x1
  Entry point address: 0x4f0
  Start of program headers: 64 (bytes into file)
  Start of section headers: 6408 (bytes into file)
  Flags:      0x0
  Size of this header: 64 (bytes)
  Size of program headers: 56 (bytes)
  Number of program headers: 9
  Size of section headers: 64 (bytes)
  Number of section headers: 28
  Section header string table index: 27
hilgerbj@os1:~/cse278/Lab9$
```

## f. ELF sections

To see a list of all the sections, use the following command:

**readelf -SW a.out**

This will result in something like the following output:

```
[hilgerbj@os1:~/cse278/Lab9$ readelf -SW a.out
```

There are 29 section headers, starting at offset 0x1b28:

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	0000000000000000	000000	000000	00		0	0	0
[ 1]	.interp	PROGBITS	0000000000000238	000238	00001c	00	A	0	0	1
[ 2]	.note.ABI-tag	NOTE	0000000000000254	000254	000020	00	A	0	0	4
[ 3]	.note.gnu.build-id	NOTE	0000000000000274	000274	000024	00	A	0	0	4
[ 4]	.gnu.hash	GNU_HASH	0000000000000298	000298	000024	00	A	5	0	8
[ 5]	.dynsym	DYNSYM	00000000000002c0	0002c0	000108	18	A	6	1	8
[ 6]	.dynstr	STRTAB	00000000000003c8	0003c8	000117	00	A	0	0	1
[ 7]	.gnu.version	VERSYM	00000000000004e0	0004e0	000016	02	A	5	0	2
[ 8]	.gnu.version_r	VERNEED	00000000000004f8	0004f8	000040	00	A	6	2	8
[ 9]	.rela.dyn	RELA	0000000000000538	000538	000108	18	A	5	0	8
[10]	.rela.plt	RELA	0000000000000640	000640	000048	18	AI	5	22	8
[11]	.init	PROGBITS	0000000000000688	000688	000017	00	AX	0	0	4
[12]	.plt	PROGBITS	00000000000006a0	0006a0	000040	10	AX	0	0	16
[13]	.plt.got	PROGBITS	00000000000006e0	0006e0	000008	08	AX	0	0	8
[14]	.text	PROGBITS	00000000000006f0	0006f0	000212	00	AX	0	0	16
[15]	.fini	PROGBITS	0000000000000904	000904	000009	00	AX	0	0	4
[16]	.rodata	PROGBITS	0000000000000910	000910	000018	00	A	0	0	4
[17]	.eh_frame_hdr	PROGBITS	0000000000000928	000928	000054	00	A	0	0	4
[18]	.eh_frame	PROGBITS	0000000000000980	000980	000168	00	A	0	0	8
[19]	.init_array	INIT_ARRAY	0000000000200d88	000d88	000010	08	WA	0	0	8
[20]	.fini_array	FINI_ARRAY	0000000000200d98	000d98	000008	08	WA	0	0	8
[21]	.dynamic	DYNAMIC	0000000000200da0	000da0	000200	10	WA	6	0	8
[22]	.got	PROGBITS	0000000000200fa0	000fa0	000060	08	WA	0	0	8
[23]	.data	PROGBITS	0000000000201000	001000	000010	00	WA	0	0	8
[24]	.bss	NOBITS	0000000000201020	001010	000118	00	WA	0	0	32
[25]	.comment	PROGBITS	0000000000000000	001010	00002b	01	MS	0	0	1
[26]	.symtab	SYMTAB	0000000000000000	001040	0006c0	18		27	47	8
[27]	.strtab	STRTAB	0000000000000000	001700	000325	00		0	0	1
[28]	.shstrtab	STRTAB	0000000000000000	001a25	0000fe	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),  
 L (link order), O (extra OS processing required), G (group), T (TLS),  
 C (compressed), x (unknown), o (OS specific), E (exclude),  
 l (large), p (processor specific)



```
Select ahmede@os1: ~
ahmede@os1:~$ readelf -SW a.out
There are 28 section headers, starting at offset 0x1928:

Section Headers:
[Nr] Name                Type              Address           Off    Size  ES Flg Lk Inf Al
[ 0]                      NULL              0000000000000000 000000 000000 00  0  0  0
[ 1] .interp                PROGBITS          000000000000238 000238 00001c 00  A  0  0  1
[ 2] .note.ABI-tag          NOTE              000000000000254 000254 000020 00  A  0  0  4
[ 3] .note.gnu.build-id     NOTE              000000000000274 000274 000024 00  A  0  0  4
[ 4] .gnu.hash              GNU_HASH          000000000000298 000298 00001c 00  A  5  0  8
[ 5] .dynsym                DYNSYM            0000000000002b8 0002b8 000090 18  A  6  1  8
[ 6] .dynstr                STRTAB            000000000000348 000348 00007d 00  A  0  0  1
[ 7] .gnu.version            VERSYM            0000000000003c6 0003c6 00000c 02  A  5  0  2
[ 8] .gnu.version_r         VERNEED           0000000000003d8 0003d8 000020 00  A  6  1  8
[ 9] .rela.dyn              RELA              0000000000003f8 0003f8 0000c0 18  A  5  0  8
[10] .init                  PROGBITS          0000000000004b8 0004b8 000017 00  AX  0  0  4
[11] .plt                   PROGBITS          0000000000004d0 0004d0 000010 10  AX  0  0 16
[12] .plt.got               PROGBITS          0000000000004e0 0004e0 000008 08  AX  0  0  8
[13] .text                  PROGBITS          0000000000004f0 0004f0 0001d2 00  AX  0  0 16
[14] .fini                  PROGBITS          0000000000006c4 0006c4 000009 00  AX  0  0  4
[15] .rodata                PROGBITS          0000000000006d0 0006d0 000004 04  AM  0  0  4
[16] .eh_frame_hdr          PROGBITS          0000000000006d4 0006d4 00004c 00  A  0  0  4
[17] .eh_frame              PROGBITS          000000000000720 000720 000148 00  A  0  0  8
[18] .init_array             INIT_ARRAY        000000000000df0 000df0 000008 08  WA  0  0  8
[19] .fini_array             FINI_ARRAY        000000000000df8 000df8 000008 08  WA  0  0  8
[20] .dynamic                DYNAMIC           000000000000e00 000e00 0001c0 10  WA  6  0  8
[21] .got                    PROGBITS          000000000000fc0 000fc0 000040 08  WA  0  0  8
[22] .data                  PROGBITS          000000000001000 001000 000010 00  WA  0  0  8
[23] .bss                   NOBITS            000000000001010 001010 000008 00  WA  0  0  1
[24] .comment                PROGBITS          000000000001010 001010 00002b 01  MS  0  0  1
[25] .symtab                SYMTAB            000000000001040 001040 0005e8 18  26 42  8
[26] .strtab                STRTAB            000000000001060 001060 000206 00  0  0  1
[27] .shstrtab              STRTAB            00000000000108e 00108e 0000f9 00  0  0  1

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
l (large), p (processor specific)

ahmede@os1:~$
```

g. What you get from the above command, explain the following sections:

- eh\_frame/.eh\_frame\_hdr

The eh\_frame and .eh\_frame\_hdr are in charge of handling exceptions and provide tables in how to unwind the stack. According to the report above, it is the PROGBITS type, which means it's a section that contains either initialized data and instructions or instructions only.

.eh\_frame\_hdr has a memory location of 0000000000000928

offset: 000928

size: 000054 (84 bits)

.eh\_frame has a memory location of 0000000000000980

Offset: 000980

Size: 000168 (360 bits)

0000000000000980 000980 000168

- .init\_array/.fini\_array/.init/.fini

The .init\_array/.fini\_array/.init/.fini sections are in charge of the initialization and termination functionality. According to the report, the .init\_array is of type INIT\_ARRAY, which means that this section contains an array of pointers to initialization functions. Then the .fini\_array is of type FINI\_ARRAY, which means that this section



contains an array of pointers to termination functions. The `.init` and `.fini` are both of type `PROGBITS`, which means they only contained already initialized data and instructions or instructions only.

`.init_array` has a memory location of `0000000000200d88`

Offset: `000d88`

Size: `000010`

`.fini_array` has a memory location of `0000000000200d98` Offset: `000d98`

Size: `000008`

`.init` has a memory location of `0000000000000688`

Offset: `000688`

Size: `000017`

`.fini` has a memory location of `0000000000000904`

Offset: `000904`

Size: `000000`

- `.dynsym`

The `.dynsym` section contains the symbol table necessary to support dynamic linking. According to the report, the `.dynsym` has a memory location of `00000000000002c0`

Offset: `0002c0`

Size: `000108`

3. Recompile the `Lab9.cpp`, this time using `-v` flag  
**g++ Lab9.cpp -v**

From the report generated, identify the **PIE** (Position Independent Executable)

A PIE is an executable that can be run from anywhere in the primary memory without worrying about the absolute memory address. This means that now that the file has been configured with PIE as shown by the flags in the report, it can be run anywhere in the primary memory and still function properly.

There are some notable sections that should be pointed out:

- `.text` (contains most of the code associated with the program)

- .data (contains global variables initialized other than 0)
  - .bss (global variables that should be initialized to 0 )
4. a. Now, modify the Lab9.cpp code, just to print one line of text as: “The answer is: 42\n”. This time don’t call any of the test() function from main().

Recompile the code

- b. Issue the following command:

**hexdump -C a.out | grep “The” -B1 -A1**

Attach the screen shot what you see from above command.

```
[hilgerbj@os1:~/cse278/Lab9$ hexdump -C a.out | grep "The" -B1 -A1
00000900  f3 c3 00 00 48 83 ec 08  48 83 c4 08 c3 00 00 00  |....H...H.....|
00000910  01 00 02 00 00 54 68 65  20 61 6e 73 77 65 72 20  |.....The answer |
00000920  69 73 3a 20 34 32 0a 00  01 1b 03 3b 54 00 00 00  |is: 42.....;T...|
hilgerbj@os1:~/cse278/Lab9$
```

- c. Answer the following questions

- What is a pipe in Linux? How we represent pipe?

We represent the pipe with the “|” character and a pipe is a form of redirecting data from one command to another. In the instance above, we are redirecting the output from the hexdump command to the grep command, which then the grep filters the data and then displays that to the console.

- Elaborate the meaning of the command grep.

Grep stands for the Global regular expression print and searches files for a specified pattern of characters and then returns the memory addresses that contain that specified pattern.

- Make an educated guess what B and A might refer to in the above command.

The B1 flag displays the contents of one (hence B1) memory address found after the main address found using grep

The A1 flag displays the contents of one (hence A1) memory address after the main address found using grep

- Use different B and A values for the above, then attach another screen shot what you get

```
[hilgerbj@os1:~/cse278/Lab9$ hexdump -C a.out | grep "The" -B5 -A3
000008c0 ff 48 85 ed 74 20 31 db 0f 1f 84 00 00 00 00 |.H..t 1.....
000008d0 4c 89 fa 4c 89 f6 44 89 ef 41 ff 14 dc 48 83 c3 |L..L..D..A..H.
000008e0 01 48 39 dd 75 ea 48 83 c4 08 5b 5d 41 5c 41 5d |.H9.u.H...[ ]A\A
000008f0 41 5e 41 5f c3 90 66 2e 0f 1f 84 00 00 00 00 |A^A_..f.....
00000900 f3 c3 00 00 48 83 ec 08 48 83 c4 08 c3 00 00 00 |...H...H.....
00000910 01 00 02 00 00 54 68 65 20 61 6e 73 77 65 72 20 |....The answer
00000920 69 73 3a 20 34 32 0a 00 01 1b 03 3b 54 00 00 00 |is: 42.....;T..
00000930 09 00 00 00 78 fd ff ff a0 00 00 00 b8 fd ff ff |...x.....
00000940 c8 00 00 00 c8 fd ff ff 70 00 00 00 d2 fe ff ff |.....p.....
[hilgerbj@os1:~/cse278/Lab9$
```

5. Issue the following command to see the loadable components:

**readelf -IW a.out**

Attach the screen shot what you see from above command

```
[hilgerbj@os1:~/cse278/Lab9$ readelf -IW a.out
Elf file type is DYN (Shared object file)
Entry point 0x0
There are 9 program headers, starting at offset 64

Program Headers:
Type           Offset  VirtAddr  PhysAddr  FileSiz  MemSiz  Flg Align
PHDR           0x000040 0x00000000 0x00000000 0x000040 0x0001f8 R  0x5
INTERP        0x000238 0x00000000 0x00000000 0x000238 0x00001c R  0x1
              [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
LOAD          0x000000 0x00000000 0x00000000 0x000000 0x000000 R E 0x200000
LOAD          0x000c00 0x00000000 0x00000000 0x000200 0x000200 R  0x200000
DYNAMIC       0x000da0 0x00000000 0x00000000 0x000da0 0x000200 R  0x8
NOTE         0x000254 0x00000000 0x00000000 0x000254 0x000044 R  0x4
GNU_EH_FRAME 0x000928 0x00000000 0x00000000 0x000928 0x000054 R  0x4
GNU_STACK    0x000000 0x00000000 0x00000000 0x000000 0x000000 R  0x10
GNU_RELRO    0x000d08 0x00000000 0x00000000 0x000d08 0x000278 R  0x1

Section to Segment mapping:
Segment Sections...
00
01 .interp
02 .interp.note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt .init .plt .plt.got .text .fini .rodata .eh_frame_hdr .eh_frame
03 .init_array .fini_array .dynamic .got .data .bss
04 .dynamic
05 .note.ABI-tag .note.gnu.build-id
06 .eh_frame_hdr
07
08 .init_array .fini_array .dynamic .got
```

## PART B: Debugger (gdb)

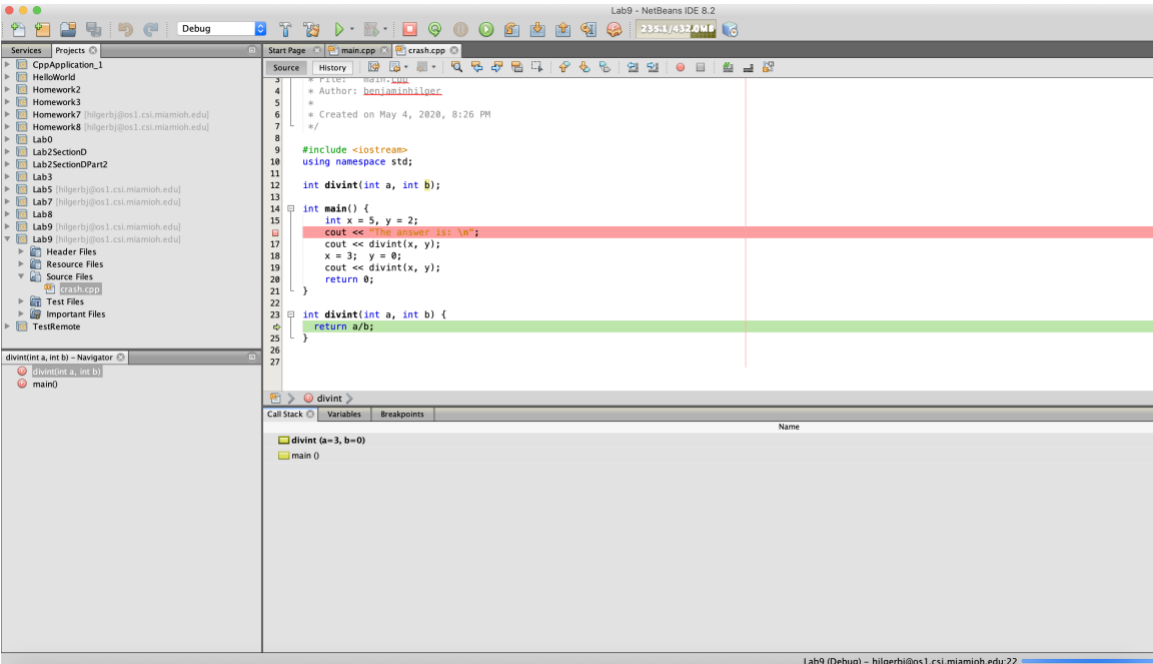
### Goals

- How to write more efficient code
- Where to look when hard-to-find bugs arise.

**Note:** You may use the supplied **crash.cpp** in [os1.csi.miamioh.edu](https://os1.csi.miamioh.edu) (compile using **-g switch**) to experiment with all the commands that are available in **gdb**, alternatively, you may use your own sample code to experiment and should use **NetBeans8.2** built-in **gdb** debugger. Here are the operations you should try:

1. Try setting and removing breakpoints
2. Step line-by-line in a method (**may be the main()**) and observe changes in a variable
3. Try navigating the stack frames in the debugger (**observe how the call stack looks**)

4. Finally, make a screenshot of your whole NetBeans window showing:
  - a. The current line of code
  - b. The stack trace tab
  - c. The variables tab



- Starting gdb
  - gdb nameOfExecutable
- Running a Program
  - (gdb) run
- Stack Trace
  - (gdb) backtrace
- Examining Variables
  - (gdb) print x
- Listing the Program
  - (gdb) list
- Setting Breakpoints
  - (gdb) help breakpoint

From command shell, issue

```
$ gdb -help
```

It will show all the commands usage

The same help can be found when running gdb already:

```
$gdb
```

```
(gdb) help
```

To quit from debugger, use the command q

```
(gdb) q
```

Command	Purpose
q /quit	Quit from gdb
L	List
l myfunction	
b main	Puts a breakpoint at the main
break 5	Set a breakpoint at line 5 and run the program
Run	
b N	Puts breakpoint at line N
r /run	Start script
Where	
Up	
p x	Print variable
Cont	Continue
commnads	End with a line saying just “end”

## Submission

- No late assignments will be accepted!
- This work is to be done individually
- The submission file will be saved with the name **Lab9\_yourMUID.pdf**
- Assignment is due by Mon/Tue May 4/5 during Lab time.
- On or before the due time, drop the *electronic copy* of your work in the *canvas*
- **Don't forget to Turn in the files! Lab9\_yourMUID.pdf & Lab9\_yourMUID\*.cpp**